

WARSAW MODEL TRAINERS

# Can a small model pass the matura?

25–27 SEP 2026  
Kolektyw3, Koszykowa 54



# Hackathon rules

## Model & training

Open model, up to 8 GB. Any model that fits under the weight limit. Can use multiple ones as long as all are under 8GB each.

## Sources

Any source, including synthetic data. As much as you want. All of internet is fine too.

## Offline only

Your harness is not allowed to connect to internet during the test.

## Your work

Your code and harness are yours. Share your work with the judges.

## Harness

You can use your own or public harnesses.

## Repository

No copyrighted content in the repo. List sources + a script to fetch them.

# Sunday exam

- 01 Training ends at 11:00.  
Final; exam will be provided then.
- 02 The model answers offline,  
without human help.
- 03 No ChatGPT, Claude, Gemini  
or other closed APIs.
- 04 Local RAG and tools are allowed.  
No model over 8 GB.
- 05 Submit the final exam results for base models  
and trained model (if you want to participate in  
the improvement track).



# Sunday submission

Full details tomorrow!



# Scoring

## 01 Best matura score

Highest final exam result

## 02 Best progress

Biggest improvement over the untouched base model(s)

## 03 Mały, ale wariat

Smallest model to pass with a score of at least 35%



# Prize pool

# PLN 8,000

A shared prize pool for teams winning in different categories

Plus Quesma rewards

**K3 Viber Passes for all winners**

Desk access at Kolektyw3



\$200 per participant for the hackathon  
Winners: an extra \$200 per member afterwards





# \$200

in Cursor credits  
for every participant



**Scan to redeem**

[Your \\$200 in Cursor credits](#)



# Mentors, experts & jury



**Tomasz Kolinko**

**Mentor**



**Ania Olchowik**

**Mentor / Jury**




**Jacek Migdał**

**Mentor**



**Michał Warda**

**Mentor**



**Piotr Migdał**

**Expert / Jury**



**Gleb Badzeika**

**Expert / Jury**



**Anna Kubasiak**

**Expert**



**Bartek Kolasa**

**Expert**



**Olaf Tomalka**

**Jury**

# At Kolektyw3

FIND US

**Koszykowa 54**

3rd floor · left from the lift

OPEN

**09:00–00:00**

**Coffee, tea & snacks**

DOOR CODE

**24568#**



# Timetable

## PROGRAMME

## FOOD

### FRIDAY

25 SEP

15:00 Workshop & check-in  
18:00 Opening  
20:25 Team building  
21:05 Work starts

20:00 Pizza

### SATURDAY

26 SEP

12:00–18:00 Mentors  
12:00 Team deadline

09:00 Breakfast  
14:30 Lunch  
20:00 Pizza

### SUNDAY

27 SEP

10:45 Presentation order  
11:00 Exams & demos  
13:00 Panel: Smok, JetBrains & Quesma  
14:00 Results & prizes

09:00 Breakfast  
12:00 Lunch

**Closing party by Smok × Bad Ideas Fund · until 22:00**

# Sponsors



**JETBRAINS**



**BADideas  
.fund**



**NEBIUS**

 **TechTree**

 **Kolektyw3**    **× AI Tinkerers**

# Compute



**Labqoat**

**Ask a mentor how to set up**

Optional: Modal free Starter plan, \$30/month compute



# Are you planning your next move?

TechTree, our partner and an AI-powered recruitment platform helping startups, scale-ups, and enterprise companies hire top talent faster  
**is here to connect you with top companies.**

Teams they hired for:

||ElevenLabs

Σ zeta labs



Kick



LocalGlobe

||ElevenLabs

Σ zeta labs



Kick



LocalGlobe

 hopin

Personio

Scan the link to discover most exciting opportunities:





Questions & contact

# See you on Telegram.



**Join the community**

Organisers, sponsors  
and mentors

[t.me/warsawmodeltrainers](https://t.me/warsawmodeltrainers)

# Register your team

## TODAY

### 01 Register once

One person creates the team and saves the TEAM\_KEY. We show it only once.

### 02 Bring your teammates

Click “Join a team”, paste the key and add your name.

### 03 Add your model later

Base model is optional now.  
Add it under “Update team”.

**One key for your team and the exam script.**



**SCAN TO REGISTER**

[warsawmodeltrainers.dev/  
matura.html#register](https://warsawmodeltrainers.dev/matura.html#register)

# Finding your team

## How many people can join a team?

There is no team-size limit.

## Can I join two teams?

No. One person, one team, one submission.

## What if I do not have a team?

Mingle.

# Team registration

## How do we register?

Go to [warsawmodeltrainers.dev](https://warsawmodeltrainers.dev).

One person registers and saves the TEAM\_KEY, shown only once. Others select “Join a team”, paste the key and enter their name.

## When is the deadline? What about the model?

Enter the full team roster by Saturday, 12:00.

The base model is optional now. Add it later under “Update team”.

## What is the TEAM\_KEY for? What if I lose it?

The same key runs your exam script. Keep it within your team. If you lose it, an organiser can retrieve it from the admin panel.

# On site & support

## Do I have to stay all weekend?

No. At least one teammate must attend Sunday's exam and presentation. The space opens 09:00–24:00. Overnight training runs remotely.

## Where do I ask for help?

Telegram: [t.me/warsawmodeltrainers](https://t.me/warsawmodeltrainers). Organisers, sponsors, compute support, mentors, announcements and help with the app are all there.

## How do I get a mentor?

Approach a mentor on Saturday. There are no advance assignments.

# The challenge

## What are we building?

An open model up to 8 GB on disk that scores well on the Polish history matura. Fine-tune it, add a harness (RAG, prompts, agents, tools), or combine both.

## Do I need to know how to train a model?

No. A harness alone can compete. Friday's 15:00 workshop covers this.

## Why do we submit the base model too?

We measure both the final exam score and improvement over the bare model. If you use several models, the baseline is the one that scores highest alone.

# Models & size limits

## Which base model can I use?

Any downloadable open weights, including quantized Bielik, PLLuM, Qwen, Gemma or Llama. Declare its name, link and size in GB in the app.

## What does the 8 GB limit mean?

The weights as run before the exam must fit in 8 GB on disk.  
A 12B model quantized to 7 GB is fine. The RAG knowledge base is excluded.

## Several models? A LoRA adapter?

Each model must fit in 8 GB separately; there is no combined-model limit.  
LoRA adapter does not count.

# Harness & external APIs

## Can I use an existing harness or framework?

Yes. Use your own or public code, libraries, agents and RAG frameworks. System prompts, few-shot examples and chain-of-thought are harness too.

## Can I call external APIs while building?

Yes, including closed LLM APIs for synthetic training data.

## What changes during the exam?

No external APIs or web search. Use your own models, each within 8 GB, and a local RAG knowledge base. Everything must answer without human help.

# Benchmarks & prior work

## Can I start from a model already tuned on history?

Yes, as a public open base model. Its existing score becomes your benchmark. Your own model tuned before the hackathon cannot pass as a bare model.

## Can I change my base model?

Yes. Sunday uses the benchmark of the model that sits the exam. Without a benchmark, you cannot compete for best progress.

## Can I use earlier work?

Public libraries and your existing tools are allowed. Graded work must be created during the hackathon, from Friday 18:00. Submitting prior fine-tuning as the bare model means disqualification.

# Submission checklist

## Which models and files do we submit?

Full team roster by Saturday, 12:00. Submit the bare base model and frozen final weights (or weights + adapters), plus the exact harness used on stage.

## What belongs in the repository?

Training and harness code, data scripts, sources and a reproducible README. Use a public GitHub repo, or give organisers and jury read access before Sunday, 11:00. A public or openly licensed repo is optional.

## What must **SOURCE.md** say?

“Made during the Warsaw Model Trainers hackathon,  
Kolektyw3, 25–27.09.2026”

This file is required to win any prize.

# Data & sources

## Where do I get data?

Find it yourself. Any legally accessible source is fine, including synthetic data from any LLM. Your team is responsible for its sources.

## Can I bypass paywalls, logins or site restrictions?

No. Do not bypass protections or download where site terms forbid it.

## Can I put other people's content in our repo?

Do not republish copyrighted scans or unlicensed datasets.  
List source links and provide a download/processing script instead.  
Organisers may require removal or exclude projects that break these rules.

# Ownership & licences

## Who owns our work?

Your team owns its code, harness, prompts and tools, and chooses their licence. Models/data built from others' content are for education and research only, unless the source licences allow more. Weights follow the base-model licence.

## Can I download Wikipedia?

Yes. Wikipedia is CC BY-SA; derived datasets must use CC BY-SA or a compatible licence. Your RAG knowledge base does not count toward 8 GB.

## Will the organisers publish our model?

Only with your team's consent and where source licences allow it. Winning models and datasets may then be shared on Hugging Face.

# Compute & credits

## Where do I get a GPU?

Nebius and Labqoat provide compute credits to every team.  
Instructions: Friday's opening and the repo. Your own laptop is fine too.  
Modal offers a free plan with \$30/month of compute.

## How do I use Labqoat? Can I train overnight?

Log in at [app.forgehand.app](https://app.forgehand.app), add your team and get a machine.  
Yes, train remotely overnight. The venue closes 24:00–09:00.

## What about Cursor credits?

Every participant gets \$200 for the hackathon.  
Use the redemption QR on the Cursor slide, or ask an organiser.

# Running the exam

## What hardware do we use?

Your own laptop or cloud machine. There is no shared evaluation hardware. The exam script queries your model and sends only the answers.

## Can my model run on a cloud GPU?

Yes, on your own or a sponsor's machine. The ban concerns external AI APIs and internet knowledge, not the physical location of your model.

## Are there time, RAM, CPU or VRAM limits?

No extra hardware limits. Each model must fit in 8 GB on disk; no external APIs. The practical limit is stage time: a few minutes for the full exam and presentation.